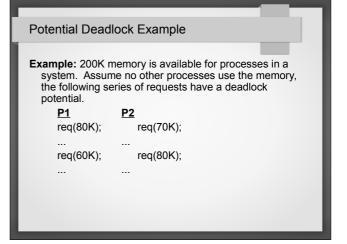
Deadlock Computer Operating Systems BLG 312E 2015-2016 Spring



Deadlock

- processes which share resources or communicate with each other become permanently blocked -> deadlock
- if processes request resources without releasing the resources they hold, deadlock may occur

Blocking Mode x Non-blocking Mode

if a resource is unavailable when requested:

- in blocking mode, process is blocked until resource becomes available
- in non-blocking mode, process receives an error message and tries later

```
Potential Deadlock Example
                               P2
   <u>P1</u>
   req(D);
                               req(T);
                               lock(T);
   lock(D);
                                 req(D);
     req(T);
                                 lock(D);
     lock(T);
                                 <....>
      <....>
                                 unlock(D);
     unlock(T);
                               unlock(T);
   unlock(D);
             Deadlock potential!
```

Potential Deadlock Example Example: If receive_msg works in blocking mode, then the following scenario has a deadlock potential. P1 P2 receive_msg(P2); receive msg(P1); send_msg(P2); send_msg (P1);

Conditions for Deadlock

- mutual exclusion condition
 - only one process can use a shared resource at a time
- hold and wait condition
 - processes wait for a requested resource until it becomes available while holding onto its own resources
- no pre-emption condition
 - resources allocated to a process cannot be taken back without the process' consent
- circular wait condition
 - two or more processes wait for the other's resource while not releasing its own in a circular fashion

Strategies used for Dealing with Deadlock

- prevention: structure the system in such a way that one of the deadlock conditions is negated
- detection and recovery: let deadlocks occur, detect them and take action
- · avoidance:
 - · don't start processes whose requests may cause a deadlock
 - · don't grant requests which may cause a deadlock
- ignore

Graph Representation

- a graph representation may be used
 - nodes in graph:

 - circle: process
 square: resource
 - edges in graph:

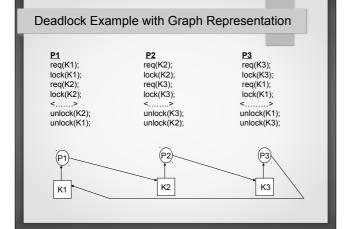
 - process → resource : process requests resource resource → process : resource allocated to process



Deadlock with circular wait

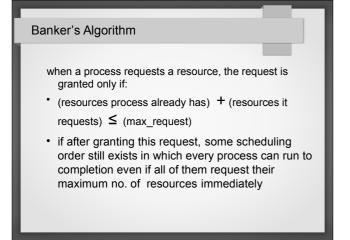
Deadlock Avoidance

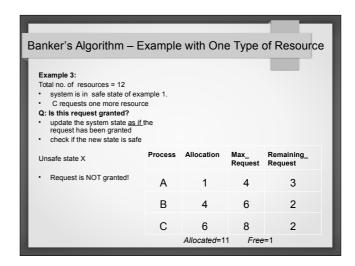
- · the Banker's algorithm
 - Dijkstra, 1965
 - fixed no. of processes and resources in the system
 - system state: current allocation of resources to processes
 - state: consists of resource and free vectors, allocation and max_request matrices

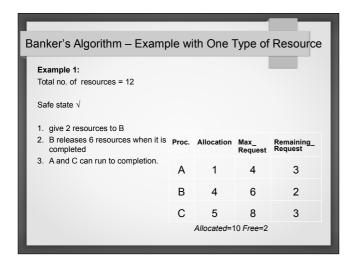


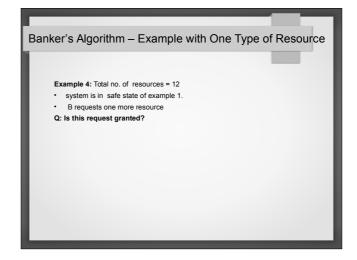
Banker's Algorithm - Definitions

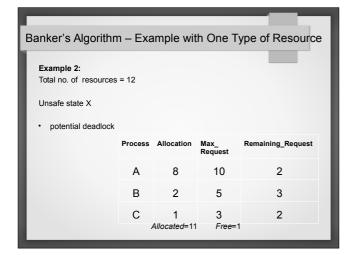
- resource: shows all resources in system
- · free: shows all free resources in system
- allocation: shows the amount of each type of resource allocated to each process
- max request: shows the maximum number of requests a process will make during its lifetime for each type of
- safe state: a state is safe if it is not deadlocked and there exists some scheduling order in which every process can run to completion even if all of them request their maximum no. of resources immediately.
- unsafe state: such a scheduling order cannot be found











Banker's Algorithm - Multiple Type of Resources Example Max_Request Matrix Allocated Matrix Remaining_Request Matrix K1 K2 K3 K1 K2 K3 0 0 P1 2 2 2 1 2 P2 6 P2 6 1 3 P2 0 0 1 P3 3 1 4 P3 2 1 1 2 2 P4 0 0 2 K1 K2 K3 0 1 1 9 3 6 if processes are executed in the order P2, P3, (P1 or P4) all may run to Resource Vector Free Vector completion \Rightarrow SAFE STATE $\sqrt{\ }$

Evaluation of the Banker's Algorithm

- to be able to apply the algorithm:
 - all processes must declare all their resource requests when they start execution
 - number of resources and processes must be fixed
 - order of process execution should not be important
 - any process holding a resource should not exit without releasing all its resources
- the algorithm grants or rejects requests based on the worst case scnearionot all rejected requests would cause a deadlock → inefficient use of resources
- the algorithm is executed each time a request is made → high cost

Banker's Algorithm - Multiple Type of Resources Example

- Q: When the system is in the safe state given in the previous slide, if P3 requests one more K3, will this request be granted?
- A: Granting this request will cause an unsafe state, so it will NOT be granted.

Deadlock Detection

- · not as restrictive as avoidance strategies
- · all requests are granted
- · system is checked for deadlock periodically
 - if deadlock is detected:
 - · terminate all deadlocked processes
 - or terminate processes one by one until deadlock is removed
 - or
- has lower cost since it is not executed on each request
- provides more efficient resource use
- period for checking for deadlock is set based on the frequency of deadlock on the system

Application of the Banker's Algorithm

 Are there any rows in the remaining_request matrix ≤ free vector?

if not: unsafe state

- Assume that the process corresponding to the row chosen above, requests all the resources it needs and finishes.
- Mark the process as completed and add all its resources to the free vector
- Repeat steps 1 and 2 until either all processes are marked as "completed" (safe state) or until a deadlock occurs (unsafe state)

Deadlock Detection Application

- · Allocation matrix and Free vector used.
- Q Request matrix defined. q_{ij} shows the amount of j type of resources process i requests
- algorithm determines processes which are not deadlocked and marks them
- · initially all processes are unmarked

Deadlock Detection Steps

- Step 1: Mark all processes which correspond to rows with all 0's in the Allocation matrix
- Step 2: Create a temporary W vector to represent the Free vector
- Step 3: Find an i for which all corresponding values in the Q matrix are LE those in the W vector (P_i must be unmarked).

 $Q_{ik} \leq W_k$, $1 \leq k \leq m$

Deadlock Detection Example																		
	R1	R2	R3	R4	R.5			R1	R2	R3	R4	R5		R1	R2	R3	R4	R5
P1	0	1	0	0	1		P1	1	0	1	1	0		2	1	1	2	1
P2	0	0	1	0	1		P2	1	1	0	0	0						
Р3	0	0	0	0	1		Р3	0	0	0	1	0			R	esour	ce Ve	ctor
P4	1	0	1	0	1		P4	0	0	0	0	0		R1	R2	R3	R4	R5
	Re	quest Matrix Q Allocation Matrix A																
2)	W : P3	's re	0 C	est	ĽΕ۱	<i>N</i> ⇒ 0) =				`					Av	ailab	ole Ve	ector
4)	No		er s			ocess					unc	l⇒	Ter	min	ate			
P1	an	d P	2 r	ema	ain	unma	ark	ced	⇒ (dea	dloc	ked	!!					

Deadlock Detection Steps

- Step 4: Terminate algorithm if no such row exists
- Step 5: If such a row exists, mark the ith process and add the corresponding row in the Allocation matrix to the W vector

 $W_k = W_k + A_{ik}$, $1 \le k \le m$

· Step6: Return to step 3.

After Deadlock Detection

- terminate all deadlocked processes
- roll-back all deadlocked processes to a previous control point in time and resume from there
 - same deadlock may occur again
- terminate deadlocked processes one by one until deadlock no longer exists
- remove allocated resources from deadlocked processes one by one until deadlock no longer exists
- •

Deadlock Detection Application

- when algorithm terminates, if there are unmarked processes ⇒ Deadlock exists
 - unmarked processes are deadlocked
- algorithm only detectes if a deadlock exists in the current state or not

Deadlocked Process Selection for Termination

- · select the one which has used the least amount of CPU
- select the one which has the longest expected time to completion
- select the one which has the least no of allocated resources
- · select the one with the lowest priority
- ...