

Memory Management Part 2

Computer Operating Systems
BLG 312E

2015-2016 Spring

Virtual Memory

- to run, a process must be in memory
 - Question: must the whole of the process be in memory?
- physical addresses are determined after a process is loaded onto the memory
 - physical addresses may be different during the whole lifetime of the process
- parts of a process don't have to be placed at contiguous locations in memory

Virtual Memory

- unused parts are in secondary memory
- initially, a part of the process is loaded onto the main memory
 - resident set
- if the part that is being accessed is not in memory
 - page fault – interrupt occurs
 - process is blocked
 - the requested part is loaded onto memory
 - operating system generates I/O request
 - interrupt occurs when I/O is completed; waiting processes are awakened and become READY

Virtual Memory

- due to virtual memory, there can be more processes in READY mode
 - more efficient multi-programming
 - only necessary parts of process are in main memory
 - processes larger than the whole main memory can also be run
- paging/segmentation is used in implementation
 - requires hardware support

Virtual Memory

Questions to answer:

- how is space allocated on the main memory and secondary storage
 - easier with paging
 - harder with segmentation due to unequal segment sizes
- what to consider when moving pages/segments between main memory \leftrightarrow secondary storage ?
- if main memory is full, which page/segment should be removed to secondary storage ?

Allocation of Memory for Unequal Sized Segments

- keep free spaces in a linked list in increasing order of their address values
- in each record of the linked list:
 - address of free space
 - size of free space
 - pointer to next free space
- add memory locations to list as they are freed
 - combine with previous and next records if possible
- de-fragmentation is useful


- **first-fit**
 - starting from the beginning of the list, allocate the first free space whose size is greater than or equal to the required size
 - leftover spaces are again added to the list
- **next-fit**
 - start looking for the first appropriate free space starting from the location of memory space allocated in the previous request (not from the beginning of the list)
 - better to have a circular list
- **best-fit**
 - try to find the free space whose size fits the requested size best (minimum leftover free space)
 - for each time, go through the whole list
- **worst-fit**
 - opposite of best-fit
 - again go through the whole list for each request

- order the free spaces in increasing order of their sizes :
 - best fit = first fit
 - harder to combine neighbor free spaces
- or keep pointers to locations in the list of free spaces of different sizes
 - takes time to update the pointers

- “buddy” system
 - divide the whole memory into blocks of size 2^k
 - assume the whole memory size is 2^s
 - there are $(s+1)$ linked lists
 - $2^0, 2^1, 2^2, \dots, 2^s$
 - list(k): pointer to blocks of size 2^k ($k=0,1,\dots,s$)
 - initially list(s) points to the first location of the memory
 - all other lists are initially empty

- assume a block of size 2^k is requested
[$>2^{k-1}$ and $\leq 2^k$]
 - if list(k) is empty, try list(k+1)
 - if not empty, split the block into two
 - add one of the resulting blocks to list(k)
 - use the other one for the request
 - if all lists are empty, the request cannot be satisfied
- when allocated blocks are returned, they are added to appropriate lists
 - “buddy” blocks are combined

Buddy system example



Fetching Techniques

- which criteria should be used when moving pages from secondary storage \Rightarrow main memory ?
 - pre-paging
 - pages that will be accessed in the near future can be predicted
 - load pages onto memory before the actual access request
 - lesser page faults
 - high costs for wrong predictions
 - good for data pages for example
 - demand paging
 - bring pages to main memory only when they are accessed

Page Replacement

- if there is no available free space in the main memory, a page needs to be moved to the secondary storage
 - care must be given to possible page traffic
 - a page that is just removed from the main memory should not be accessed
 - "thrashing": loss of time
 - main aim is to NOT remove USEFUL pages
 - pages that won't be accessed in the near future can be removed
 - some operating system pages cannot be removed
 - frame locking is done through setting a bit
 - page selection can be at two levels :
 - local: choose from among the pages of the running process
 - global: choose from among all the pages

Page Replacement

- select randomly
 - easy to implement
 - USEFUL pages may be selected
- first in first out – FIFO
 - select page which has been in the main memory the longest
 - performance may be bad – the oldest page may not be the page that won't be accessed in the near future
- BIFO (biased FIFO)
 - select from among the n_i pages of the i . process, use FIFO for the n_i pages
 - Different processes may have different number of pages in memory
 - n_i for each process may change over time

Page Replacement

- LRU (Least Recently Used)
 - high implementation cost, hardware support needed
 - keep a table of records for each page of the time that has passed since the last access to that page
 - at the end of each quantum, all entries are updated
 - clear the access time counters for the accessed pages
 - increment the access time counters for all other pages in the main memory (the ones that were not accessed)
 - when choosing a page to remove from memory, choose the one with the highest counter value (means the page has not been accessed for the longest time)

Page Replacement – Example

Assume there are 4 page frames in memory and the following pages are accessed in the given order. Initially all page frames are empty.

1 – 2 – 3 – 4 – 1 – 5 – 3 – 6 – 7 – 5 – 2

Give the contents of the page frames after each access. Mark the page faults.

- (a) Use LRU
(b) Use FIFO

Page Replacement

- use pre-defined priorities
 - the compiler can determine which page should have higher priority (do not remove from memory)
 - the structure of the program can give this info
 - principle of locality): tendency of code and data accessed to remain in the same area
 - e.g. loops, data lists

Page Replacement

- use system defined priorities
 - possible to use the same priorities used in scheduling
 - in case of a page fault, a page belonging to the process with the lowest priority is selected to be removed from main memory
 - e.g. through LRU
 - if the last page of the process with the lowest priority is removed, it has to wait until space becomes available in the main memory
 - PROBLEM: there may be unused pages of higher priority processes in main memory

Page Replacement

- use hybrid techniques
 - some rules can be combined
 - in order of decreasing preference:
 - select a read-only-access page of a blocked process
 - select a read/write-access page of a blocked process
 - select an operating system page that has not been access during the previous $\frac{1}{2}$ seconds
 - select a page of a process waiting for I/O
 - select a page of an active (running)